

We thank the reviewers for the time and effort invested in the detailed reviews and the valuable suggestions and comments to improve the quality of the paper. Below we reply to the individual points of the reviewers.

1 Rebuttal to Review 1

There should, however, be some discussion somewhere in the paper on the relationship between path length and KB size. Even if a KB uses a non-parallelizable ontology language this doesn't matter much if the paths that cause problems are short compared to the size of the KB. The non-parallel aspects of path processing of short paths will not be noticeable with a small number of threads (certainly up to hundreds) and maybe not even for very many threads (because of general overhead).

We added xxx to adress this...

The second sentence of Section 3 is badly stated and misleading. It is definitely true that for a Datalog program $\langle R, I \rangle$ the rule set R is fixed. However, the initial facts I are also fixed.

We made it clear that the rule set is fixed for a class of Datalog programs. We also revised the similar presentations in Section 2.

It is not clear whether the classes of Datalog programs considered in the paper have to share the same rule set. The discussion previous to Definition 1 seems to so indicate, but there is nothing in the definition that says so.

We modified Definition 1 to make it clear that the Datalog programs in a class share the same rule set.

There is no definition of what the parents of a node in a graph are. However, it is quite obvious what the definition should be.

We added the definition of the parents of a node in a materialization graph in Definition 2.

Example 6 does need a bit of revision. It is not the case that $Oex3$ cannot be materialized tractably in parallel. Instead there has to be a class of ontologies that cannot be so materialized.

We did the revision for Example 6 (see the paragraph following Example 6), and Example 1 (see the last paragraph of Section 3.2).

It would be much better if the ontology names corresponded to the example numbers, i.e., the ontology in Example 6 was $Oex6$.

We made the ontology subscripts correspond to the examples numbers for Example 5, Example 6, Example 7 and Example 9.

The notion of a simple concept is very much tied to the algorithm developed in the paper. How often do actual ontologies have constructs that could make them intractable that are rendered tractable by having simple concepts? If there are none, then this class is uninteresting, particularly as the paper makes such strong claims about worrying about particular ontologies. This is different from an ontology only containing simple concepts.

...

The paper is all about parallel tractability of materialization, both theoretical and practical. However the paper gives insufficient attention to RDFS, which is the main ontology language used in practice, and particularly for large knowledge bases. Even though RDFS is laughably simple it does not have parallel tractability of even entailment. The paper should better examine parallel tractability in the context of RDFS or RDFS with minor extensions.

Of course, RDFS does not fit into the approach of the paper, as it does not a-priori separate classes and individuals. The paper needs to discuss this issue, as it prominently mentions RDFS.

...

The paper does a bad job of describing why materialization in RDFS is not parallelizable. RDFS KBs have an infinite number of consequences, so it is certainly not possible to materialize all of them quickly. There are, however, well known tricks that can be used to generate a finite representation of this infinite set of consequences. However, even this cannot be done quickly in parallel. In fact, even RDFS entailment cannot be done quickly in parallel, because RDFS allows for a kind of twisted interaction (which I think is similar to twisted paths in the paper) when computing subproperty relationships. Of course, having subproperties of `rdfs:subPropertyOf` is not a usual thing to do in RDFS. The paper needs to be much more clear on its relationship to RDFS.

...

The paper mentions YAGO early on but does not indicate why YAGO admits good performance for parallel reasoning. YAGO uses a very simple ontology language - RDFS plus two simple extensions. It is thus a prime candidate for close examination to show how practical results differ from worst-case theoretical ones. However, the paper does not have any examination of reasoning on YAGO. Such an

examination is very much called for. Similarly an examination of the performance of other large RDFS-based KBs, such as DBpedia, would be very useful. If YAGO is not benchmarked then it needs to be dropped entirely. Similarly for LUBM.

...

The benchmarking part of the paper is very limited and extremely hard to follow. Only eight ontologies were benchmarked. Why were these eight chosen? Only considering at most eight threads is much too constricting.

...ontologies must satisfy the properties for parallel tractability...

More detail is needed on how the KBs were constructed. Chain length is a very important aspect of the KBs, but nothing is said of how the instances in the KB are connected together. It is not stated which ontologies belong to which DL.

...

Figure 5 is very misleading because it uses different scales for the two systems. There are several immediate questions that arise from this Figure that are poorly handled in the paper. First, why does going from one thread to two make such a difference for ParallelDHL? Without a convincing explanation for this surprising aspect of the benchmarking it is hard to trust any of the rest of the benchmarking. The other surprising aspect of Figure 5 is that RDFox speeds up so little. There is a short explanation of why this is so, but I would have liked to see a longer one.

...

The computation of the speedup numbers is wholly unexplained. Why should this number have any importance?

...

There is a claim that the benchmarking distinguishes between the Ddhlo and D-dhlo. However, Figure 5 doesn't really show this well, if at all. For example, why should processing slow down with more threads at all? Further, it appears that even for several of the more well-behaved ontologies that there is very little speedup between 6 and 8 processors.

...

The benchmarking should be redone with more ontologies and with more threads. As well, there needs to be better analysis of where parallelism is failing. For example, if chains are short, then there should be no problem in achieving near-perfect thread utilization, ignoring memory contention problems.

Wording and structural problems:

The paper has so many grammatical errors near its beginning that I gave up marking them after the first two pages. I give a few that I noticed here.

"parallelly tractable" is very grating. It would be better to rewrite sentences to not use it.

TODO: Remove the quote above once the issues are fixed.

The paper has been carefully proof-read and hopefully all wording and structural issues are fixed now.

The abstract of the paper is much too long. An abstract is supposed to be one not-long paragraph. Instead it is here two long paragraphs. The abstract needs to be cut down to a suitable size before any publication.

TODO: shorten the abstract.

Indeed, the abstract has been shortened significantly now.

2 Rebuttal to Review 2

First, in the discussion under Lemma 1, it says $|P^|$ is polynomial in the size of P , which is imprecise, as $|P^*|$ is polynomial in the size of I but not necessarily in the size of R . This is not a big issue if one considers NC for data complexity. What is less straightforward and more of concern is the statement that Step 2 costs only one time unit. As Step 2 involves checking the applicability of $B_1, \dots, B_n \rightarrow H$ in G and updating G , it depends on the size of G and the size of G is not a constant w.r.t. that of I . Footnote 4 suggests checking applicability of the rule can be done in constant time via an index. What about updating G and maintaining the index?*

In the paragraph under Lemma 1, we made it clear that $|P^*|$ is polynomial in the size of I for a class of datalog programs.

The statement that Step 2 costs only one time unit is based on the context of one processor. Since one processor is allocated one rule instance (see the first paragraph of Algorithm A_{bsc}), the procedure can be completed in constant time units. We rewrote these sentences by clarifying this.

The constant time of checking applicability by a processor is based on a common assumption when analyzing the parallel complexity. That is the problem is encoded on a parallel configuration. Consider materializing datalog programs by Algorithm A_{bsc} . This assumption allows us mapping all possible substitutions of the atoms and all rule instances in P^* to a unique memory location. Since each rule instance is allocated to a processor, there is a one-to-one relation between processors and rule instances. In this way, a processor can check the applicability of its corresponding rule instance and access the states of an atom occurring in this rule instance in constant time. Let a , b and c represent the maximum arity of any EDB predicate, the maximum number of variables in any datalog rule, and the number of datalog rules respectively. Since a , b and c depend only on the rule set, they are constants. We further have that the number of constant is at most $|\mathbf{I}|a$ and the number of all possible rule instances in P^* is at most $c(|\mathbf{I}|a)^b$. Thus, the memory space for storing atoms and rule instances is polynomial in the size of \mathbf{I} . It is like that there is a polynomial size of index that maps each atom and rule instance to the corresponding memory location.

The above assumption is also used for analyzing the computation of transitive closures and the problem of boolean matrix multiplication. In the revision, we gave detailed information about the assumption.

The computation of the rch relation S_{rch} (\dagger) looks problematic, as it requires “H has been added to G”. If it were the case, then in Example 4 the S_{rch} would always be empty, as none of the rule heads H would be added to G . Again, the complexity analysis of A_{opt} is unclear and possibly flawed. Step 2 uses an NC algorithm to compute the transitive closure of S_{rch} , which runs in poly-logarithmic time w.r.t. the size of S_{rch} not necessarily so w.r.t. the size of I . What is the size of S_{rch} w.r.t. that of I ? Similarly as above, it involves updating G , which depends on the sizes of both G and S_{rch} . Would it be still in poly-logarithmic time?

The requirement “H has been added to G” in the description of computing the rch relation S_{rch} (\dagger) is indeed incorrect. It should be “H has not been added to G”. We fixed this error.

(TODO) We added a paragraph to discuss the complexity of Algorithm A_{opt} . In Step 2 of Algorithm A_{opt} , an NC algorithm is used to compute the transitive closure of S_{rch} . It can be checked that the scale of S_{rch} is less than that of P^* and is in polynomial of the size of inputs. Thus, the NC complexity of the computation of the transitive closure can be guaranteed.

Finally, the algorithms assume each processor handling a rule instantiation. I wonder how it is implemented in ParallelDHL.

(TODO) We used a paragraph to discuss this after Algorithm A_{prc} .

I list some relatively minor issues regarding presentation and technical details:

- p20, it mentions a group of materialisation rules in YAGO which belong to D_{dhl} . It deserves a clarification on what these rules are like and why they do not fall into existing parallel tractable classes.
- p20, Table 2, it should mention how these ontologies were selected.
- p21, the discussion on line graph lg1 is confusing (i.e. “line-2 stays higher than line-6, but lower than line-4”). Maybe I was looking at the wrong place.
- p25, the section heading “Reference” occur twice.

TODO: remove the quote once the issues are fixed.

All minor issues have been fixed.